

Viraj Dudhat — Today's Interview Questions

Only today's interview questions from the stored-procedure section onward.

1. Were they forcing the data through a stored procedure?
2. So the PMS/DB team had written the stored procedure to extract the data and from there you got the data?
3. There are five different sources, right?
4. Did you connect to only one server for client information, or different PMS providers/systems?
5. So the count is only one? You extracted from the same SQL Server only?
6. So from only one SQL Server, you were creating one pipeline for getting Client information?
7. When data comes into Raw, what format does it come in?
8. Does it come as JSON, CSV, or another format?
9. How many records were coming daily in production?
10. Were records coming in one file or multiple files?
11. If hundreds of files come in a day, how do you read those files?
12. How do you read the files using Spark?
13. What do you mean by file name/file path when reading?
14. Suppose JSON structure changes, e.g. age string → integer. How do you handle it?
15. If tomorrow data comes as integer, how do you handle it?
16. What different types of compute are available in Databricks?
17. Which compute have you used?
18. Why did you use that compute?
19. Serverless is provisioned on the fly and has less spin-up time than a normal dedicated cluster, right?
20. How do you handle schema evolution in Spark?
21. You said JSON data from Client; how did you handle JSON arrays and child objects?
22. How do you share data from Delta Lake to external systems?
23. If an application needs to access data from Delta Lake, how do you provide access from Databricks?
24. What configuration is required to connect Power BI to Databricks?
25. When Power BI queries Databricks, how is the Databricks cost calculated?
26. Job compute/serverless has a cost; when Power BI runs queries, how does the Databricks-side cost work?
27. Have you heard of RLS (Row-Level Security)?
28. How would you implement RLS if a manager should see only limited users/team?

29. If you need a KPI for the last four years, how would you design/implement it?
30. Using Python, if you want to consume data from a REST API like ADP/ServiceNow/mainframe, how would you consume it?
31. Did you work hands-on consuming data from REST APIs?
32. Did you have any questions for us?