

DataNimbus Mock Interview - Questions Asked

Databricks / PySpark / Delta Lake / Project Discussion

- 1 Why are you pulling client information?
- 2 Why is another stored procedure needed? Why can't the main stored procedure or regular SQL handle the data?
- 3 If some required source data is missing, do you still proceed to the next layer?
- 4 How do you integrate stored procedures into Databricks?
- 5 How are you consuming data returned by a stored procedure?
- 6 Why are you using PyODBC instead of Spark JDBC?
- 7 What is the size of the data returned by a stored procedure?
- 8 How do you handle large volumes of data while pulling from stored procedures?
- 9 If a DataFrame has a struct-type column, how do you convert each struct field into a separate column?
- 10 What is an objective function and how does it work?
- 11 What is F.lit() in PySpark?
- 12 What is the small file problem in data engineering?
- 13 What is predicate pushdown?
- 14 What are the different types of clusters available in Databricks?
- 15 A table was initially performing well, but after many inserts, updates, and deletes, queries became slow. What could be the reasons and fixes?
- 16 What is Liquid Clustering?
- 17 What are the constraints or limitations of Liquid Clustering?
- 18 If the data is in TBs and you have a Boolean column, will clustering/indexing on that column still require a large scan?
- 19 When does shuffle happen in Spark?
- 20 How do you prevent or reduce shuffle?
- 21 What is a Concurrent Modification Exception in Delta Lake?
- 22 How do you handle a Concurrent Modification Exception?
- 23 How do you prevent multiple jobs from accessing/updating the same Delta table at the same time?
- 24 What if 10 different jobs are trying to update the same Delta table?
- 25 How can we ensure data quality in a Databricks Delta pipeline?
- 26 How is data quality handled in Delta Live Tables?
- 27 What is Lakeflow Declarative Pipelines?
- 28 Which things are not supported on Databricks Serverless compute?
- 29 Suppose a table has an array-type column and you use explode(). Why might some records disappear?
- 30 What is the solution when records disappear after using explode()?
- 31 How can I check the history of a Delta table?
- 32 If I update a Delta table, how is that update processed internally by Delta Lake?
- 33 How does a Delta table process an UPDATE internally?
- 34 Have you worked with or explored Genie, Genie One, Genie Code, or Genie Agents?
- 35 How can Genie be used with AI/BI dashboards or Power BI dashboards?

- 36 If we consume Change Data Feed from a table and an update happens, what records are produced?
- 37 If one employee record changes from age 25 to 26, how many CDF records are generated?
- 38 What additional metadata columns are provided by Delta Change Data Feed?

Total Questions: 38